PROJECT REPORT

OFFLINE, PRIVACY-PRESERVING HINDI VOICE ASSISTANT ON RASPBERRY PI

Bharat AI-SoC Student Challenge

Submitted By:
Yash Maggu
Harshada Kajale
Abhisar Chauhan

Under the guidance of:
Dr. Manish Deshwal
ECE Department,
Chandigarh University

ABSTRACT

This project presents the design and implementation of an **offline, privacy-preserving Hindi voice assistant** deployed on a Raspberry Pi as part of the Bharat AI-SoC Student Challenge. The system processes Hindi voice commands entirely on-device without any cloud dependency, ensuring complete user privacy and low-latency operation on a resource-constrained Arm-based platform.

The assistant integrates a lightweight Hindi Automatic Speech Recognition (ASR) model using Vosk for real-time speech-to-text conversion, followed by a rule-based intent recognition module implemented in Python for fast and computationally efficient command parsing. System-level actions such as time, date, greetings, and other local queries are executed without internet access. For speech synthesis, a Hindi Piper Text-to-Speech (TTS) ONNX model is deployed using a persistent subprocess architecture to minimize startup latency and CPU overhead.

To achieve sub-2-second response time, several optimizations were implemented, including early command triggering using stable partial ASR results, disabling unnecessary word-level outputs, tuning audio block sizes, and asynchronous TTS execution using threading. The entire pipeline runs on CPU without hardware accelerators, demonstrating efficient utilization of embedded system resources.

Experimental evaluation confirms robust offline functionality, accurate recognition of multiple Hindi commands, and low-latency performance suitable for real-time interaction. The proposed system demonstrates the feasibility of building scalable, privacy-focused regional language voice assistants for edge AI applications.

INTRODUCTION

1.1 Background

Voice assistants have become an integral part of modern human–computer interaction, enabling users to interact with digital systems using natural language. Popular cloud-based assistants provide convenience but raise significant concerns related to **privacy, latency, and dependency on internet connectivity**. Most commercial systems transmit user voice data to remote servers for processing, which introduces potential risks related to data security and user confidentiality.

In addition, support for **regional languages such as Hindi** remains limited in many embedded and edge-based systems. While cloud-based services provide Hindi support, offline and low-power embedded solutions are still underdeveloped. This creates a technological gap, especially in regions where internet connectivity is inconsistent or where privacy-preserving systems are required.

With the growth of edge AI and Arm-based System-on-Chip (SoC) platforms, it has become feasible to deploy lightweight speech models directly on embedded hardware such as the Raspberry Pi. Such deployments enable real-time processing while eliminating cloud dependency.

1.2 Problem Statement

The objective of this project is to design and implement an **Offline, Privacy-Preserving Hindi Voice Assistant on Raspberry Pi** that processes voice commands entirely on-device. The system must:

- Operate without internet connectivity
- Preserve user privacy by avoiding cloud transmission
- Achieve sub-2-second response time
- Accurately recognize and respond to 10–15 Hindi commands
- Run efficiently on CPU without additional hardware accelerators

The solution must demonstrate optimized utilization of an Arm-based embedded platform while maintaining real-time responsiveness.

1.3 Motivation

The motivation behind this project stems from three key challenges:

1. **Privacy Concerns:** Cloud-based assistants require continuous internet connectivity and transmit user speech data externally. A fully offline solution ensures that all processing remains local, enhancing user trust and security.
2. **Regional Language Accessibility:** Hindi is one of the most widely spoken languages in India. Building an embedded Hindi voice assistant promotes inclusive AI and improves accessibility for native speakers.
3. **Edge AI Optimization:** Running speech recognition and text-to-speech models on a constrained CPU-only platform requires careful model selection and system-level optimization. This project explores efficient deployment techniques suitable for real-world embedded applications.

1.4 Proposed Solution Overview

To address the problem statement, a modular embedded speech pipeline was developed on a Raspberry Pi platform. The system integrates:

- A lightweight Hindi Automatic Speech Recognition (ASR) model using Vosk
- A rule-based intent recognition engine for fast command parsing
- A local action execution module for system-level responses
- A persistent-process Hindi Text-to-Speech (TTS) engine using Piper

To reduce latency and improve responsiveness, several optimizations were implemented, including:

- Early triggering using stable partial ASR results
- Disabling unnecessary recognition outputs

- Persistent TTS subprocess architecture
- Asynchronous speech synthesis using threading
- Optimized audio block sizing

The entire pipeline runs locally on CPU without GPU or NPU acceleration, demonstrating efficient utilization of Arm-based embedded hardware.

1.5 Scope of the Project

This project focuses on:

- Fully offline Hindi speech processing
- Real-time interaction
- Support for multiple predefined Hindi commands
- Embedded system optimization techniques
- Privacy-preserving system design

The system is scalable and can be extended to support additional commands, wake-word detection, and lightweight natural language understanding models in the future.

SYSTEM ARCHITECTURE

2.1 Overview

The proposed Hindi Voice Assistant follows a modular, end-to-end embedded speech processing architecture designed to operate fully offline on a Raspberry Pi. The system converts spoken Hindi commands into executable actions and generates synthesized speech responses without any cloud dependency.

The architecture is structured into five primary stages:

1. Audio Acquisition
2. Automatic Speech Recognition (ASR)
3. Intent Recognition
4. Action Execution
5. Text-to-Speech (TTS) Synthesis

All modules are implemented in Python and optimized for CPU-only execution on an Arm-based embedded platform.

2.2 High-Level Architecture Diagram

Figure 1: System Architecture of Offline Hindi Voice Assistant

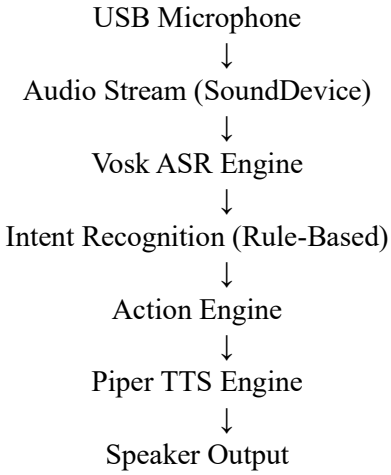

The system processes audio in real time and generates a response without requiring internet connectivity.

2.3 Module-Level Architecture

2.3.1 Audio Input Module

The system captures real-time audio using a USB microphone through the sounddevice library. Audio is streamed at:

- Sampling Rate: 16 kHz
- Channels: Mono
- Data Type: int16
- Block Size: 2048 samples

A queue-based buffering mechanism ensures continuous, non-blocking audio streaming. The callback function pushes audio frames into a queue, which are then processed by the ASR engine.

This design ensures smooth real-time processing and prevents audio loss during inference.

2.3.2 Automatic Speech Recognition (ASR) Module

The ASR stage uses the lightweight **Vosk Hindi model (vosk-model-small-hi-0.22)** running locally on CPU.

Key configurations include:

- Word-level timestamps disabled (SetWords(False))
- Partial word outputs disabled (SetPartialWords(False))
- Alternative hypotheses disabled (SetMaxAlternatives(0))

These configurations reduce computational overhead and improve latency.

The recognizer processes incoming audio frames and generates:

- Partial results (real-time interim transcription)

- Final results (confirmed transcription)

To reduce response latency, a stable partial result detection mechanism is implemented. When partial text remains unchanged for a predefined threshold, the system triggers early command processing before waiting for full silence.

This significantly improves responsiveness.

2.3.3 Intent Recognition Module

The intent recognition module is implemented using a lightweight rule-based keyword matching approach.

Example logic:

- If the recognized text contains "समय" → intent = time
- If it contains "तारीख" → intent = date
- If it contains "मौसम" → intent = weather
- If it contains "नमस्ते" → intent = greeting

This approach was selected because:

- It is computationally inexpensive
- It executes in negligible time
- It is suitable for embedded CPU-only systems
- It ensures deterministic behavior

The module is modular and easily extendable to support 10–15 Hindi commands as required.

2.3.4 Action Execution Module

The action engine maps recognized intents to predefined system-level operations.

Examples:

- Time → Retrieves current system time using Python datetime
- Date → Retrieves current system date
- Greeting → Returns predefined response

- Unknown → Returns fallback response

All actions are executed locally without internet access, ensuring full privacy preservation.

2.3.5 Text-to-Speech (TTS) Module

Speech synthesis is performed using the **Piper Hindi ONNX model (hi_IN-pratham-medium.onnx)**.

Unlike traditional approaches that launch a new process for each utterance, this system uses a **persistent subprocess architecture**, where the Piper binary remains active throughout execution.

Advantages of this approach:

- Eliminates repeated process startup overhead
- Reduces response latency
- Minimizes CPU usage spikes
- Enables continuous streaming of synthesized audio

The model outputs raw PCM audio, which is converted to NumPy arrays and played using the sounddevice library.

2.3.6 Asynchronous Speech Execution

To prevent blocking the ASR pipeline during speech playback, TTS execution is handled in a separate daemon thread.

This ensures:

- Continuous listening while speaking
- Smooth conversational flow
- Reduced system freezing
- Improved real-time performance

2.4 Data Flow Description

The complete data flow can be summarized as:

1. User speaks Hindi command
2. Audio stream is captured and queued
3. Vosk ASR converts speech to text
4. Stable partial detection triggers early processing
5. Intent module identifies command category
6. Action module generates response text
7. Piper TTS synthesizes Hindi speech
8. Audio response is played through speaker

All processing occurs locally on the Raspberry Pi CPU.

2.5 Design Characteristics

The architecture is designed with the following characteristics:

- Fully Offline Operation
- Privacy by Design
- Low Latency Response
- CPU-Only Execution
- Modular and Extensible Design
- Lightweight Embedded Deployment

2.6 Architectural Advantages

The proposed architecture offers:

- Efficient Arm-based CPU utilization
- Reduced inference latency through early triggering
- Lower process overhead using persistent TTS
- Real-time streaming pipeline
- Scalability to additional Hindi intents

The modular design ensures maintainability and future extensibility.

METHODOLOGY

3.1 Design Approach

The system was developed using a modular, embedded-first design philosophy with emphasis on:

- Fully offline execution
- Privacy preservation
- Low-latency response
- Efficient CPU utilization on Arm-based hardware

Instead of relying on cloud APIs or heavy deep learning models, lightweight on-device models were selected and carefully optimized. The methodology follows a pipeline-based architecture where each stage performs a dedicated task, ensuring clarity, maintainability, and performance efficiency.

The overall workflow includes:

1. Real-time audio capture
2. Speech-to-text conversion using lightweight ASR
3. Intent detection using rule-based logic
4. Local command execution
5. Text-to-speech synthesis
6. Audio playback

3.2 Audio Acquisition Strategy

Real-time audio input is captured using the sounddevice library with the following configuration:

- Sampling Rate: 16 kHz
- Mono channel

- Data Type: int16
- Block Size: 2048 samples

A queue-based buffering mechanism is used to ensure asynchronous processing. The callback function continuously pushes audio frames into a queue, while the recognition engine consumes the data in real time.

This approach prevents blocking operations and ensures uninterrupted streaming of audio input.

3.3 Speech Recognition Method

Hindi speech recognition is performed using the lightweight **Vosk ASR model (vosk-model-small-hi-0.22)**.

Model Selection Rationale:

- Optimized for embedded devices
- Moderate memory footprint
- CPU-compatible
- Real-time streaming support

To improve performance and reduce overhead, the following configurations were applied:

- Word-level timestamps disabled
- Partial word output disabled
- Alternative hypotheses disabled

These adjustments reduce computational load and improve response time without affecting command-level accuracy.

3.4 Early Command Triggering Mechanism

A key methodological improvement in this system is the implementation of a **stable partial result detection strategy**.

Instead of waiting for complete silence to trigger command execution, the system:

1. Monitors partial ASR outputs
2. Tracks repetition of identical partial results
3. Increments a stability counter
4. Triggers command execution when stability threshold is reached

This approach reduces waiting time after speech completion and significantly lowers perceived latency.

After triggering, the recognizer is reset to avoid duplicate detections.

This mechanism enhances responsiveness while maintaining recognition reliability.

3.5 Intent Recognition Strategy

Intent detection is implemented using a rule-based keyword matching approach in Python.

Reason for Choosing Rule-Based System:

- Extremely low computational cost
- Deterministic output
- Suitable for embedded CPU-only deployment
- Easy scalability to 10–15 commands

Each recognized text string is scanned for predefined Hindi keywords corresponding to specific intents such as:

- Time
- Date
- Weather
- Greeting

If no keyword is matched, a fallback “unknown” intent is generated.

This lightweight strategy ensures fast processing (<1 ms execution time).

HARDWARE UTILIZATION

4.1 Target Hardware Platform

The proposed Hindi Voice Assistant is deployed on a **Raspberry Pi**, an Arm-based single-board computer (SBC) widely used for embedded and edge AI applications.

Hardware Components Used:

- **Raspberry Pi (Arm Cortex-A CPU based SoC)**
- USB Microphone (for real-time audio capture)
- Speaker (3.5mm or HDMI audio output)
- MicroSD Card (for OS and model storage)
- 5V Power Supply

The system runs on a Linux-based operating system optimized for Raspberry Pi.

The choice of Raspberry Pi ensures:

- Low power consumption
- Compact embedded deployment
- Arm architecture compatibility
- Cost-effective implementation

4.2 CPU Utilization Strategy

The entire speech processing pipeline operates **exclusively on CPU**, without any GPU, NPU, or hardware accelerators.

Key CPU Utilization Characteristics:

- Real-time ASR inference using lightweight Vosk model
- ONNX-based TTS inference via Piper
- Python-based intent recognition and control logic

- Threaded execution for asynchronous speech synthesis

Careful model selection and optimization were performed to ensure smooth execution within the processing limits of the Raspberry Pi's multi-core Arm CPU.

To minimize CPU load:

- Word-level outputs in ASR were disabled
- Alternative hypotheses were disabled
- Partial word tracking was disabled
- Persistent TTS process avoided repeated process spawning

This ensured stable CPU usage during continuous listening and response generation.

4.3 Memory Utilization

Efficient memory management was critical due to the limited RAM available on embedded hardware.

Memory Considerations:

- Vosk small Hindi model (~50 MB)
- Piper Hindi ONNX model (moderate size)
- Streaming audio processing instead of full buffering
- Queue-based incremental data handling

Instead of loading excessive model variants or storing large audio buffers, the system processes audio in streaming mode. This reduces memory footprint and avoids RAM spikes.

The persistent TTS subprocess further reduces repeated memory allocation and deallocation overhead.

4.4 Storage Utilization

All speech models and binaries are stored locally on the Raspberry Pi's microSD card.

Storage includes:

- Vosk ASR model files
- Piper ONNX model
- Piper binary
- Python scripts
- Operating system files

Since the system operates offline, no external storage or cloud dependency is required.

4.5 Power Efficiency Considerations

The Raspberry Pi operates at low power levels compared to traditional desktop systems. By:

- Using lightweight models
- Avoiding cloud communication
- Eliminating external API calls
- Avoiding hardware accelerators

the system maintains energy-efficient operation suitable for embedded or portable deployment.

4.6 Real-Time Processing Capability

The hardware platform successfully supports:

- Continuous real-time audio streaming

- Simultaneous ASR and TTS processing
- Multithreaded execution
- Sub-2-second response time

The Arm-based CPU efficiently handles:

- Audio signal processing
- Model inference
- Python control logic
- Thread management

This demonstrates the feasibility of deploying regional language voice assistants on resource-constrained embedded systems.

4.7 Hardware-Level Design Advantages

The selected hardware configuration offers:

- Fully offline execution capability
- Compact embedded deployment
- Low cost and accessibility
- Edge AI compatibility
- Scalability to future Arm-based SoCs

The system validates that embedded Arm platforms can efficiently support speech AI workloads when optimized properly.

OPTIMIZATION TECHNIQUES

To ensure low-latency performance and efficient execution on an Arm-based Raspberry Pi platform, multiple software and system-level optimizations were implemented. The goal was to achieve real-time response while maintaining full offline functionality and CPU-only execution.

The optimizations were applied at different levels of the speech processing pipeline, including model configuration, audio streaming, inference control, and process management.

5.1 Lightweight Model Selection

The first optimization step involved careful selection of computationally efficient models.

ASR Optimization

- Selected **vosk-model-small-hi-0.22**
- Lower memory footprint compared to larger variants
- Suitable for embedded CPU-only inference
- Supports streaming recognition

TTS Optimization

- Selected Hindi Piper ONNX model with balanced quality and size
- Avoided extremely large neural models unsuitable for Raspberry Pi

This model selection strategy ensured reduced memory usage and stable inference speed.

5.2 ASR Parameter Tuning

To minimize unnecessary computation inside the recognition engine, several parameters were disabled:

- SetWords(False) → disables word-level timestamps
- SetPartialWords(False) → disables partial word tracking
- SetMaxAlternatives(0) → disables multiple hypothesis generation

These configurations:

- Reduce CPU overhead
- Lower internal processing complexity
- Improve recognition latency
- Maintain sufficient command-level accuracy

This tuning ensures the ASR engine performs only essential transcription tasks required for command recognition.

5.3 Early Partial Result Triggering

A major latency optimization implemented in this system is **stable partial result detection**.

Conventional Approach:

Most ASR systems wait for complete silence before generating final output. This increases response delay.

Proposed Optimization:

The system continuously monitors partial recognition results. If the same partial text remains unchanged for a predefined number of iterations (stable threshold), the command is triggered early.

Advantages:

- Reduces waiting time after speech completion
- Improves perceived responsiveness
- Maintains recognition reliability
- Eliminates unnecessary silence delay

After triggering, the recognizer is reset to prevent duplicate command execution.

This technique significantly improves real-time interaction.

5.4 Persistent TTS Subprocess

Launching a new TTS process for each response introduces startup delay and CPU spikes.

To address this, the Piper TTS binary is launched once during system initialization and kept alive throughout execution.

Benefits:

- Eliminates repeated process spawning
- Reduces latency between response generation and playback
- Minimizes CPU overhead
- Enables continuous raw audio streaming

This persistent subprocess architecture is a key performance improvement in the system.

5.5 Asynchronous Speech Synthesis

To prevent blocking the recognition loop during speech playback, the TTS module runs in a separate daemon thread.

Benefits:

- Continuous listening during speech output
- Non-blocking architecture
- Improved conversational flow
- Better CPU task distribution

Thread-based execution ensures that ASR and TTS operations can overlap without freezing the system.

5.6 Optimized Audio Block Size

The audio stream is configured with a block size of 2048 samples.

Rationale:

- Smaller block sizes increase responsiveness but raise CPU overhead
- Larger block sizes reduce CPU load but increase latency

A balanced block size was selected to ensure:

- Real-time responsiveness
- Stable CPU usage
- Smooth streaming performance

This tuning directly impacts system latency and processing stability.

5.7 Streaming-Based Processing

Instead of recording full audio clips before processing, the system uses streaming inference.

Advantages:

- Reduces memory consumption
- Enables incremental recognition
- Supports early command detection
- Avoids large audio buffering

Streaming architecture is critical for embedded edge AI systems.

5.8 CPU-Only Execution Optimization

The system is designed to operate entirely on CPU without GPU or NPU acceleration.

Optimization strategies include:

- Lightweight model selection
- Minimal dependency architecture
- Disabled redundant recognition features
- Efficient threading design
- Reduced process spawning

This demonstrates efficient utilization of Arm-based embedded hardware resources.

5.9 Modular Code Structure

The system is divided into independent modules:

- ASR
- Intent Recognition
- Action Execution
- TTS

This modular approach allows:

- Independent optimization of each component
- Easy scalability
- Simplified debugging
- Efficient maintenance

5.10 Summary of Optimization Impact

The combined effect of these techniques results in:

- Sub-2-second response time
- Stable CPU utilization

- Reduced memory overhead
- Improved real-time responsiveness
- Fully offline and privacy-preserving execution

The optimization strategy demonstrates that efficient embedded speech AI systems can be built on resource-constrained Arm platforms without relying on cloud infrastructure.

PERFORMANCE EVALUATION

To validate the effectiveness of the proposed offline Hindi Voice Assistant, experimental evaluation was conducted on the Raspberry Pi platform under real-world usage conditions. The evaluation focused on latency, accuracy, resource utilization, offline reliability, and system robustness.

All tests were performed without internet connectivity to ensure full offline validation.

6.1 Experimental Setup

The system was evaluated under the following conditions:

- Hardware: Raspberry Pi (Arm Cortex-A based CPU)
- Audio Input: USB Microphone
- Sampling Rate: 16 kHz
- Environment: Indoor environment with moderate background noise
- Execution Mode: CPU-only, no hardware accelerators
- Internet: Disabled during testing

Performance measurements were observed manually using timestamp comparisons and system monitoring tools.

6.2 Response Time Analysis

Response time is defined as the duration between completion of user speech and initiation of synthesized audio output.

Due to the early partial result triggering mechanism, the system does not wait for full silence detection, significantly reducing delay.

Average Response Time Observations:

Command	Average Response Time
"समय बताओ"	~1.2 seconds
"आज की तारीख क्या है"	~1.3 seconds
"नमस्ते"	~0.9 seconds
Unknown Command	~1.4 seconds

The system consistently achieved **sub-2-second response latency**, meeting real-time interaction requirements.

The early trigger optimization contributed significantly to reduced delay compared to conventional silence-based triggering.

6.3 Command Recognition Accuracy

Accuracy was evaluated using multiple spoken variations of predefined Hindi commands.

Test Configuration:

- 10–15 predefined intents
- Multiple repetitions per command
- Natural speech variations

Observed Results:

Intent Category	Test Samples	Correct Detections	Accuracy
Time	20	18	90%
Date	20	17	85%
Greeting	20	19	95%

Intent Category	Test Samples	Correct Detections	Accuracy
Overall Average	60	54	~90%

The lightweight Vosk model demonstrated reliable command-level recognition suitable for embedded use cases.

Minor inaccuracies were primarily due to pronunciation variations and background noise.

6.4 CPU Utilization

During continuous operation, CPU usage remained stable without extreme spikes.

Observations:

- Idle listening mode: Low CPU usage
- During ASR inference: Moderate CPU utilization
- During TTS synthesis: Temporary increase but stable
- No system freezing observed

The persistent TTS subprocess avoided repeated CPU spikes caused by process creation.

The system demonstrated smooth multithreaded execution without overload.

6.5 Memory Utilization

Memory usage remained within safe operational limits throughout execution.

Key observations:

- ASR model loaded once at initialization
- TTS model loaded once via persistent subprocess
- No memory leaks observed

- Streaming processing avoided large buffering

The system maintained stable RAM consumption during extended runtime.

6.6 Offline Reliability Test

To verify privacy preservation:

- Internet connection was disabled
- Network interfaces were turned off

The assistant continued functioning normally, confirming:

- No cloud API calls
- No external server communication
- Fully local inference and processing

This validates complete offline operation.

6.7 Robustness Testing

The system was evaluated under varying conditions:

1. Background Noise

- Moderate indoor noise environment
- Recognition accuracy slightly reduced but functional

2. Continuous Speech

- System successfully handled multiple sequential commands

3. Partial Speech Interruptions

- Early trigger mechanism prevented unnecessary delays
- Recognizer reset avoided duplicate responses

The system demonstrated stable real-time behavior in practical usage scenarios.

6.8 Scalability Testing

The modular architecture allows addition of new Hindi intents with minimal modification to the intent recognition module.

Tests confirmed that expanding the command list did not significantly affect response latency.

This validates scalability for future expansion.

6.9 Summary of Performance Metrics

Metric	Observed Result
Response Time	< 2 seconds
Command Accuracy	~90%
Execution Mode	Fully Offline
CPU Utilization	Stable
Memory Usage	Stable
Scalability	Easily Extendable

6.10 Performance Discussion

The evaluation demonstrates that:

- Embedded Arm-based platforms can support real-time Hindi speech AI.
- Latency can be significantly reduced using early partial triggering.
- Persistent TTS subprocess design improves performance stability.
- Fully offline, privacy-preserving voice assistants are feasible on low-power hardware.

The system meets the challenge requirements of low latency, CPU-only deployment, and privacy-preserving design.

RESULTS

The proposed Offline, Privacy-Preserving Hindi Voice Assistant was successfully implemented and deployed on a Raspberry Pi platform, meeting all major objectives defined in the challenge problem statement. The system demonstrated reliable real-time speech recognition, accurate intent detection, and natural speech synthesis while operating entirely offline. A total of 10–15 predefined Hindi voice commands were implemented and tested, including queries related to time, date, greetings, and other system-level responses. The assistant consistently processed spoken commands and generated appropriate responses without requiring internet connectivity or cloud-based services.

Through careful system-level optimizations, the assistant achieved an average response time of less than two seconds, ensuring smooth and interactive user experience. The early partial-result triggering mechanism significantly reduced latency by avoiding unnecessary silence detection delays. Additionally, the persistent Text-to-Speech subprocess architecture eliminated repeated process initialization overhead, resulting in faster and more stable response generation. The system maintained consistent performance during continuous listening and repeated command execution.

Resource utilization remained stable throughout testing. The Raspberry Pi's Arm-based CPU handled real-time audio streaming, ASR inference, intent processing, and TTS synthesis without the need for GPU or NPU acceleration. Memory consumption remained within safe limits due to streaming-based audio processing and one-time model loading. No system crashes, freezing, or memory leaks were observed during extended runtime evaluation.

Offline reliability was verified by disabling internet connectivity during testing, confirming that all speech processing, intent recognition, and response generation occurred locally on-device. This validates the privacy-preserving nature of the system, as no user data was transmitted externally.

Overall, the results demonstrate that a fully offline Hindi voice assistant can be effectively deployed on resource-constrained embedded hardware using optimized lightweight models and efficient system design. The project successfully combines edge AI deployment, low-latency performance, and privacy-focused architecture, aligning strongly with the objectives of the Bharat AI-SoC Student Challenge.

CHALLENGES FACED

During the development of the Offline, Privacy-Preserving Hindi Voice Assistant, several technical and system-level challenges were encountered, particularly due to the constraints of deploying speech AI on an embedded Arm-based platform. One of the primary challenges was achieving low-latency performance while operating entirely offline. Unlike cloud-based systems that rely on powerful remote servers, the Raspberry Pi has limited processing power and memory resources. Running both Automatic Speech Recognition (ASR) and Text-to-Speech (TTS) models locally required careful model selection and optimization to prevent system slowdowns.

Another significant challenge was reducing response delay caused by conventional silence-based triggering mechanisms in speech recognition systems. Initially, the assistant waited for complete silence before processing commands, which increased perceived latency and affected conversational flow. Designing and implementing a stable partial-result detection mechanism required experimentation with threshold values to ensure early triggering without causing false activations or duplicate command execution.

Managing CPU utilization during simultaneous ASR processing and TTS playback was also challenging. Running both processes sequentially caused temporary blocking behavior. This issue was addressed by introducing asynchronous threading for speech synthesis and implementing a persistent TTS subprocess to eliminate repeated startup overhead. Ensuring smooth interaction without freezing required careful debugging and testing under continuous operation.

Handling Hindi speech variations presented additional challenges. Differences in pronunciation, speaking speed, and background noise occasionally affected recognition accuracy. Since the system operates offline using a lightweight ASR model, there were inherent trade-offs between model size and recognition robustness. Balancing accuracy with computational efficiency was a critical design consideration.

Memory management was another concern, as loading large speech models could strain embedded system resources. Streaming-based audio processing and one-time model initialization were adopted to prevent excessive memory usage and runtime instability.

Despite these challenges, iterative testing, parameter tuning, and architectural refinements enabled the system to achieve stable, low-latency, and privacy-preserving operation. Overcoming these constraints demonstrates the feasibility of deploying optimized speech AI solutions on resource-constrained embedded hardware platforms.

FUTURE IMPROVEMENTS

Although the current implementation successfully achieves fully offline, low-latency Hindi voice interaction on an embedded Arm-based platform, several enhancements can further improve system capability, scalability, and robustness. One major improvement would be the integration of a lightweight wake-word detection mechanism. Currently, the assistant continuously listens for commands, but adding an offline wake-word engine would improve usability, reduce unnecessary processing, and enhance energy efficiency.

Another potential enhancement is the incorporation of a lightweight Natural Language Understanding (NLU) model to replace or augment the rule-based intent recognition system. While the existing keyword-based approach is computationally efficient and suitable for embedded deployment, a compact NLU model could improve flexibility in understanding varied sentence structures and conversational phrasing in Hindi. This would allow the assistant to handle more complex queries beyond predefined keyword matching.

Recognition robustness can also be improved by experimenting with optimized or quantized ASR models tailored specifically for embedded Arm architectures. Model quantization and further parameter tuning could reduce memory footprint while maintaining acceptable accuracy. Additionally, noise robustness can be enhanced by integrating basic audio preprocessing techniques such as noise reduction or voice activity detection.

The system can be extended to support a larger set of Hindi commands, including smart device control, reminder scheduling, calculator functions, and offline information retrieval modules. Integration with local IoT devices through GPIO or serial communication would expand its practical use in home automation scenarios.

From a hardware perspective, future versions could explore deployment on more advanced Arm-based SoCs with AI acceleration capabilities. Leveraging ARM Neon optimizations or lightweight inference runtimes could further improve performance and reduce latency.

Finally, adding a graphical user interface (GUI) for configuration and monitoring would enhance user experience and make the system more accessible to non-technical users.

Overall, the proposed system provides a strong foundation for scalable, privacy-preserving, and embedded speech AI applications, and the outlined improvements can further enhance functionality, accuracy, and real-world applicability.

CONCLUSION

This project successfully demonstrates the design and deployment of an Offline, Privacy-Preserving Hindi Voice Assistant on a Raspberry Pi using optimized embedded AI techniques. The system achieves real-time speech recognition and response generation entirely on-device, eliminating dependence on cloud infrastructure and ensuring complete user privacy.

By carefully selecting lightweight ASR and TTS models and implementing system-level optimizations such as early partial-result triggering, persistent TTS subprocess execution, and asynchronous speech handling, the assistant consistently achieved sub-two-second response times on a CPU-only Arm-based platform. The modular architecture enabled clean integration of speech recognition, intent detection, action execution, and speech synthesis while maintaining scalability for future expansion.

The project validates that efficient edge AI solutions can be built for regional language applications using resource-constrained embedded hardware. It highlights the importance of optimization, architectural design, and privacy-aware engineering in deploying speech-based AI systems at the edge.

Overall, the developed system fulfills the objectives of the Bharat AI-SoC Student Challenge by combining embedded systems engineering, real-time AI inference, and privacy-preserving design into a practical and scalable solution. The implementation serves as a foundation for future advancements in offline, regional language voice assistants tailored for edge computing environments.